

# Capstone Project 2: *Online Retail*

## Data Science Professional Certificate/HarvardX

*Paw Hermansen*

*June 17, 2019*

### Executive Summary

#### The Data

I examine online retail sales data found in the **Online Retail** dataset that contains all transactions for about a year around 2011 for a UK based online retail. The products are unique gift items and many of the customers are wholesalers.

The dataset contains 25900 invoices for 4373 customers. It needed only little cleaning - I deleted invoices with no customer, I removed a few invoice rows with non-product items like handling and postage, and I removed two cancelled orders with exceptionally large quantities.

After cleaning, the dataset contains 21784 invoices for 4362 customers.

#### The Goals

##### Part 1: Customer Segmentation (Clustering)

RFM, **R**ecency, **F**requency, and **M**onetary Value, is an often used way to measure the value of a customer. I introduce a new Customer Segmentation based on the RFM data computed by the Kmeans clustering algorithm and compare it to a traditional fixed Customer Segmentation defined by PUTLER also based on RFM.

##### Part 2: Predicting Customer Segment for new Customers (Classification)

I examine if it is possible to predict which Customer Segment a new user will end up in at the last date of the dataset based on the size of the customers first invoice in the dataset and the customers demographics, i.e. the customers country.


I classify to both of the two kinds of customer segment sets computed in the first part, and I use and compare the K-nearest neighbor (knn) algorithm, the Random Forest algorithm, and an approach with random segment selection. I chose the knn and Random Forest algorithms because they are base on very different principles.

#### The Results


##### The Result Part 1: Customer Segmentation

Each dot in the following two plots is a customer. The position represents the R, F, and M values and the colors represent the customer segment in each of the two systems:

## Putler Customer Segments



## K-means Customer Segments



The Putler segmentation has eleven segments and is shown for comparison in the left plot. The new Customer Segmentation computed by Kmeans clustering has seven segments and is shown in the right plot. The percentages of customers in each segment and the names I gave the segments are:

	Name	Cust. pct	Description
1	Small and Occasional	9.2%	Recently active but spends low and not very often
2	New	57.8%	Fairly recently active but spend low and have not been active earlier
3	Good Average	0.9%	Active and buys from seldom to often and spends from low to average
4	Big Spenders	0.1%	Spend most money of all and they recently did it again
5	Inactive	14.1%	Not active for a long time and spend low
6	Fading Away	17.7%	Not active for some time and spend low
7	Good and Loyal	0.1%	Spend average or more money and buy very often

## The Result Part 2: Predicting Customer Segment

The accuracy of predicting the customer segment from the customers first invoice is:

Method	Putler Segments Accuracy	Kmeans Segments Accuracy
Random	0.127	0.355
K-nearest neighbor	0.137	0.306
Random Forest	0.178	0.391

## Conclusion

### Customer Segmentation

My Machine Learning generated customer segments ends up with seven customer segments compared to eleven for the PUTLER approach. Other traditional RFM approaches, however, splits the the customers into between about five and twenty segments, and so the my Machine Learning generated customer segments are within what can be considered normal. The analysis, though, shows that the two sets of customer segments splits the customers in very different ways, and the segments of the two approaches do not look like each other. It is impossible to say which set of the segment sets that are best because it depends on the intended use.

Further work could be done to make an analysis of the development over time of customers. To find dynamic trends or customer lifecycles would however require a dataset that covers a longer period of at least several years for a business like the one, I analysed in this report.

### Predicting Customer Segment

The K-nearest neighbor algorithm predicted the future customer segment worse than randomly for the new Kmeans segments! The Random Forest algorithm predicted the future customer segment of both customer segment sets a little bit better than random and also better the K-nearest neighbor algorithm. The accuracy is still low for both segment sets, though. Predicting using Random Forest might be a beginning of helping the business to a finer segmentation of new customers.

It is clearly easier to predict the Kmeans segments - this is quite naturally because of the fewer segments and the very high number of customers in a single segment.

In this report I did not use the information of the content of the first order and it would be interesting to investigate if using this data could help increase the accuracy of the predicting. Also the Random Forest analyses had a problem finding the best value of `mtry` and this should be fixed.

## The Data and the Goal

I have chosen to examine online retail sales data found in the **Online Retail** dataset from <https://archive.ics.uci.edu/ml/datasets/Online+Retail>. The documentation describes the dataset as:

This is a transnational data set which contains all the transactions occurring between 01/12/2010 and 09/12/2011 for a UK-based and registered non-store online retail. The company mainly sells unique all-occasion gifts. Many customers of the company are wholesalers.

In the first part of this report I compare a traditional RFM Customer Segmentation of this data to a new Customer Segmentation of this data computed by the Kmeans clustering algorithm.

In the second part of this report I examine if it is possible to predict which Customer Segment a new user will end up in after a year, i.e classification of new customers. I classify to both of the two kinds of customer segment sets computed in the first part, and I use and compare the K-nearest neighbor (knn) algorithm and the Random Forest algorithm for the classifications. I chose knn and Random Forest because they work on different principles.

# Data Exploration and Cleaning

## Basic Information

The original data `Online Retail.xlsx` looks like:

```
## Observations: 541,909
## Variables: 8
## $ InvoiceNo    <chr> "536365", "536365", "536365", "536365", "536365", ...
## $ StockCode   <chr> "85123A", "71053", "84406B", "84029G", "84029E", "...
## $ Description <chr> "WHITE HANGING HEART T-LIGHT HOLDER", "WHITE METAL...
## $ Quantity    <dbl> 6, 6, 8, 6, 6, 2, 6, 6, 6, 32, 6, 6, 8, 6, 6, 3, 2...
## $ InvoiceDate  <dtm> 2010-12-01 07:26:00, 2010-12-01 07:26:00, 2010-12...
## $ UnitPrice   <dbl> 2.55, 3.39, 2.75, 3.39, 3.39, 7.65, 4.25, 1.85, 1....
## $ CustomerID  <dbl> 17850, 17850, 17850, 17850, 17850, 17850, 17850, 1...
## $ Country     <chr> "United Kingdom", "United Kingdom", "United Kingdo...
```

Each of the 541909 rows is one line of an invoice. Rows with the same `InvoiceNo` make up an invoice for a customer.

The documentation of the dataset describes the fields:

- **InvoiceNo:** Invoice number. Nominal, a 6-digit integral number uniquely assigned to each transaction. If this code starts with letter 'c', it indicates a cancellation.
- **StockCode:** Product (item) code. Nominal, a 5-digit integral number uniquely assigned to each distinct product.
- **Description:** Product (item) name. Nominal.
- **Quantity:** The quantities of each product (item) per transaction. Numeric.
- **InvoiceDate:** Invoice Date and time. Numeric, the day and time when each transaction was generated.
- **UnitPrice:** Unit price. Numeric, Product price per unit in sterling.
- **CustomerID:** Customer number. Nominal, a 5-digit integral number uniquely assigned to each customer.
- **Country:** Country name. Nominal, the name of the country where each customer resides.

The data contains 25900 different invoices and 4373 different customers. The temporal extension of the data is from 2010-12-01 07:26:00 to 2011-12-09 11:50:00, i.e. a couple days more than one year of data.

The first few lines of the data and a summary of the data is:

```
## InvoiceNo StockCode Description Quantity
## 1 536365 85123A WHITE HANGING HEART T-LIGHT HOLDER 6
## 2 536365 71053 WHITE METAL LANTERN 6
## 3 536365 84406B CREAM CUPID HEARTS COAT HANGER 8
## 4 536365 84029G KNITTED UNION FLAG HOT WATER BOTTLE 6
## 5 536365 84029E RED WOOLLY HOTTIE WHITE HEART. 6
## 6 536365 22752 SET 7 BABUSHKA NESTING BOXES 2
## InvoiceDate UnitPrice CustomerID Country
## 1 2010-12-01 07:26:00 2.55 17850 United Kingdom
## 2 2010-12-01 07:26:00 3.39 17850 United Kingdom
## 3 2010-12-01 07:26:00 2.75 17850 United Kingdom
## 4 2010-12-01 07:26:00 3.39 17850 United Kingdom
## 5 2010-12-01 07:26:00 3.39 17850 United Kingdom
## 6 2010-12-01 07:26:00 7.65 17850 United Kingdom
```

```

## InvoiceNo      StockCode      Description
## Length:541909 Length:541909 Length:541909
## Class :character Class :character Class :character
## Mode :character Mode :character Mode :character
##
##
##
##
## Quantity      InvoiceDate      UnitPrice
## Min. : -80995.00 Min. :2010-12-01 07:26:00 Min. : -11062.06
## 1st Qu.: 1.00 1st Qu.:2011-03-28 10:34:00 1st Qu.: 1.25
## Median : 3.00 Median :2011-07-19 16:17:00 Median : 2.08
## Mean : 9.55 Mean :2011-07-04 13:02:44 Mean : 4.61
## 3rd Qu.: 10.00 3rd Qu.:2011-10-19 10:27:00 3rd Qu.: 4.13
## Max. : 80995.00 Max. :2011-12-09 11:50:00 Max. : 38970.00
##
## CustomerID      Country
## Min. :12346 Length:541909
## 1st Qu.:13953 Class :character
## Median :15152 Mode :character
## Mean :15288
## 3rd Qu.:16791
## Max. :18287
## NA's :135080

```

From the summary we make two suspicious observations that need to be investigated further later in this report:

- The lowest quantity is negative and the highest quantity is the positive of the same number.
- The lowest unit-price is negative and the highest unit-price is nearly UK £ 40,000.

## Explore and Clean the Data

### Missing data

The number of missing data in each field is:

```

## Rows with missing InvoiceNo = 0

## Rows with missing StockCode = 0

## Rows with missing Description = 1454

## Rows with missing Quantity = 0

## Rows with missing InvoiceDate = 0

## Rows with missing UnitPrice = 0

## Rows with missing CustomerID = 135080

```

```
## Rows with missing Country = 0
```

I want to analyse the customers and I cannot use invoice rows with no customers. But before I remove the rows with no `CustomerID` I check if all invoices have exactly one customer. The following prints the number of invoices that have rows with different customerID's or have a mix of NA's and different customerID's:

```
onlineRetail %>%
  group_by(InvoiceNo) %>%
  summarize(numberOfCustomerIDs = length(unique(CustomerID))) %>%
  filter(numberOfCustomerIDs != 1) %>%
  nrow()
```

```
## [1] 0
```

This means that no invoices have rows with a mix of customers, or no customer, and I delete all rows without `CustomerID`:

```
onlineRetail = onlineRetail[!is.na(onlineRetail$CustomerID), ]
nrow(onlineRetail)
```

```
## [1] 406829
```

And now all the empty descriptions are gone too:

```
## Description rows with missing data = 0
```

## InvoiceNo

The data contains 22190 different invoice numbers, and all invoice numbers have the right format of 6 digits with or without a 'C' prefixed, as the resulting 0 in the following code shows. Nothing needs to be cleaned here:

```
nrow(onlineRetail[!grepl("^C\\d{6}$", onlineRetail$InvoiceNo) & !grepl("^\\d{6}$", onlineRetail$InvoiceNo), ])
```

```
## [1] 0
```

## StockCode

The data contains 3684 different stock codes. The documentation says that `StockCode` consist of 5 digits. I try to print up to ten rows where `StockCode` does not consist of exactly 5 digits:

```
head(onlineRetail[!grepl("^\\d{5}$", onlineRetail$StockCode), ], 10)
```

##	InvoiceNo	StockCode	Description	Quantity
## 1	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6
## 3	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8
## 4	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6
## 5	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6
## 46	536370	POST	POSTAGE	3

```
## 50      536373      85123A WHITE HANGING HEART T-LIGHT HOLDER      6
## 52      536373      84406B      CREAM CUPID HEARTS COAT HANGER      8
## 61      536373      82494L      WOODEN FRAME ANTIQUE WHITE      6
## 62      536373      84029G KNITTED UNION FLAG HOT WATER BOTTLE      6
## 63      536373      84029E      RED WOOLLY HOTTIE WHITE HEART.      6
##              InvoiceDate UnitPrice CustomerID      Country
## 1  2010-12-01 07:26:00      2.55      17850 United Kingdom
## 3  2010-12-01 07:26:00      2.75      17850 United Kingdom
## 4  2010-12-01 07:26:00      3.39      17850 United Kingdom
## 5  2010-12-01 07:26:00      3.39      17850 United Kingdom
## 46 2010-12-01 07:45:00     18.00      12583      France
## 50 2010-12-01 08:02:00      2.55      17850 United Kingdom
## 52 2010-12-01 08:02:00      2.75      17850 United Kingdom
## 61 2010-12-01 08:02:00      2.55      17850 United Kingdom
## 62 2010-12-01 08:02:00      3.39      17850 United Kingdom
## 63 2010-12-01 08:02:00      3.39      17850 United Kingdom
```

So it is seen that not all `StockCode` follow the documentation. Analysing a little deeper we see that no `StockCode` contains a lowercase letter:

```
nrow(onlineRetail[onlineRetail$StockCode != toupper(onlineRetail$StockCode), ])
```

```
## [1] 0
```

We see, though, that the `StockCode` has a much broader format than told in the documentation. One other format is to append one or two letters to the `StockCode` which occurs in 33889 rows. Apparently this is often used to indicate the product color as in this example:

```
nrow(onlineRetail[grepl("^\\d{5}[A-Z]{1,2}$", onlineRetail$StockCode), ])
```

```
## [1] 33889
```

All these products seem to be genuine and valid products and I leave them in the data.

Product codes that differ from five digits with possible one or two letters appended, looks like:

```
head(onlineRetail[!grepl("^\\d{5}[A-Z]{0,2}$", onlineRetail$StockCode), ], 10)
```

```
##      InvoiceNo      StockCode Description Quantity      InvoiceDate
## 46      536370          POST      POSTAGE         3 2010-12-01 07:45:00
## 142     C536379           D      Discount        -1 2010-12-01 08:41:00
## 387     536403          POST      POSTAGE         1 2010-12-01 10:27:00
## 1124    536527          POST      POSTAGE         1 2010-12-01 12:04:00
## 1424    536540          C2      CARRIAGE         1 2010-12-01 13:05:00
## 2240    536569           M      Manual          1 2010-12-01 14:35:00
## 2251    536569           M      Manual          1 2010-12-01 14:35:00
## 4407    536779 BANK CHARGES Bank Charges         1 2010-12-02 14:08:00
## 5074    536840          POST      POSTAGE         1 2010-12-02 17:27:00
## 5259    536852          POST      POSTAGE         1 2010-12-03 08:51:00
##              UnitPrice CustomerID      Country
## 46          18.00      12583      France
## 142          27.50     14527 United Kingdom
```

```
## 387      15.00      12791      Netherlands
## 1124     18.00      12662          Germany
## 1424     50.00      14911             EIRE
## 2240      1.25      16274 United Kingdom
## 2251     18.95      16274 United Kingdom
## 4407     15.00      15823 United Kingdom
## 5074     18.00      12738          Germany
## 5259     18.00      12686          France
```

When collected to a full list of all the different ones we see that they all are about non product items as handling and postage:

```
onlineRetail[!grepl("^\\d{5}[A-Z]{0,2}$", onlineRetail$StockCode), ] %>%
  group_by(StockCode) %>%
  count(Description) %>%
  top_n(1, wt=n) %>%
  ungroup() %>%
  select(-n)
```

```
## # A tibble: 8 x 2
##   StockCode Description
##   <chr>      <chr>
## 1 BANK CHARGES Bank Charges
## 2 C2         CARRIAGE
## 3 CRUK       CRUK Commission
## 4 D          Discount
## 5 DOT        DOTCOM POSTAGE
## 6 M          Manual
## 7 PADS       PADS TO MATCH ALL CUSHIONS
## 8 POST       POSTAGE
```

I remove them from the data because they only are non product items. It will probably not matter much because there only are 1920 invoice rows with one of these product codes:

```
onlineRetail = onlineRetail[grepl("^\\d{5}[A-Z]{0,2}$", onlineRetail$StockCode),, ]
nrow(onlineRetail)
```

```
## [1] 404909
```

## Cancelled Invoices

Invoices with InvoiceNo starting with a 'c' is a cancellation invoice that cancels one or more rows in other invoices. For example:

```
head(onlineRetail[startsWith(tolower(onlineRetail$InvoiceNo), "c"), ], 10)
```

```
##      InvoiceNo StockCode      Description Quantity
## 155    C536383   35004C SET OF 3 COLOURED FLYING DUCKS      -1
## 236    C536391   22556  PLASTERS IN TIN CIRCUS PARADE     -12
## 237    C536391   21984  PACK OF 12 PINK PAISLEY TISSUES    -24
## 238    C536391   21983  PACK OF 12 BLUE PAISLEY TISSUES    -24
```



##	239	C536391	21980	PACK OF 12 RED RETROSPOT TISSUES	-24
##	240	C536391	21484	CHICK GREY HOT WATER BOTTLE	-12
##	241	C536391	22557	PLASTERS IN TIN VINTAGE PAISLEY	-12
##	242	C536391	22553	PLASTERS IN TIN SKULLS	-24
##	940	C536506	22960	JAM MAKING SET WITH JARS	-6
##	1442	C536543	22632	HAND WARMER RED RETROSPOT	-1
##		InvoiceDate	UnitPrice	CustomerID	Country
##	155	2010-12-01 08:49:00	4.65	15311	United Kingdom
##	236	2010-12-01 09:24:00	1.65	17548	United Kingdom
##	237	2010-12-01 09:24:00	0.29	17548	United Kingdom
##	238	2010-12-01 09:24:00	0.29	17548	United Kingdom
##	239	2010-12-01 09:24:00	0.29	17548	United Kingdom
##	240	2010-12-01 09:24:00	3.45	17548	United Kingdom
##	241	2010-12-01 09:24:00	1.65	17548	United Kingdom
##	242	2010-12-01 09:24:00	1.65	17548	United Kingdom
##	940	2010-12-01 11:38:00	4.25	17897	United Kingdom
##	1442	2010-12-01 13:30:00	2.10	17841	United Kingdom

In these examples all the quantities are negative, and it can be shown that negative quantities correspond to exactly the cancellation invoices because there are no rows with a cancellation invoice with a positive quantity and no rows with a non-cancellation invoice with a negative quantity:

```
nrow(onlineRetail[startsWith(tolower(onlineRetail$InvoiceNo), "c") ^ onlineRetail$Quantity < 0, ])
```

```
## [1] 0
```

The cancelled rows do not show which other invoice row they cancel. In some cases a probable original invoice row can be found manually, like the suspicious observation about the minimum and maximum quantity made earlier:

```
onlineRetail[abs(onlineRetail$Quantity) >= 40000, ]
```

##		InvoiceNo	StockCode	Description	Quantity
##	61620	541431	23166	MEDIUM CERAMIC TOP STORAGE JAR	74215
##	61625	C541433	23166	MEDIUM CERAMIC TOP STORAGE JAR	-74215
##	540422	581483	23843	PAPER CRAFT , LITTLE BIRDIE	80995
##	540423	C581484	23843	PAPER CRAFT , LITTLE BIRDIE	-80995
##		InvoiceDate	UnitPrice	CustomerID	Country
##	61620	2011-01-18 09:01:00	1.04	12346	United Kingdom
##	61625	2011-01-18 09:17:00	1.04	12346	United Kingdom
##	540422	2011-12-09 08:15:00	2.08	16446	United Kingdom
##	540423	2011-12-09 08:27:00	2.08	16446	United Kingdom

In other cases the cancelled original invoice row cannot be found. For example listing all the `onlineRetail` rows for customer 17548 show three invoices where one of them, C536391, is a cancellation but there is no rows showing that the customer ever ordered these items - perhaps because the original order was made before the first date of this data:

```
onlineRetail[onlineRetail$CustomerID == 17548, ]
```

##	InvoiceNo	StockCode	Description	Quantity
----	-----------	-----------	-------------	----------

##	236	C536391	22556	PLASTERS IN TIN CIRCUS PARADE	-12
##	237	C536391	21984	PACK OF 12 PINK PAISLEY TISSUES	-24
##	238	C536391	21983	PACK OF 12 BLUE PAISLEY TISSUES	-24
##	239	C536391	21980	PACK OF 12 RED RETROSPOT TISSUES	-24
##	240	C536391	21484	CHICK GREY HOT WATER BOTTLE	-12
##	241	C536391	22557	PLASTERS IN TIN VINTAGE PAISLEY	-12
##	242	C536391	22553	PLASTERS IN TIN SKULLS	-24
##	165025	550755	22585	PACK OF 6 BIRDY GIFT TAGS	24
##	165026	550755	22082	RIBBON REEL STRIPES DESIGN	10
##	165027	550755	22081	RIBBON REEL FLORA + FAUNA	10
##	165028	550755	22079	RIBBON REEL HEARTS DESIGN	10
##	165029	550755	22926	IVORY GIANT GARDEN THERMOMETER	4
##	177224	C552049	22926	IVORY GIANT GARDEN THERMOMETER	-4
##	177225	C552049	22585	PACK OF 6 BIRDY GIFT TAGS	-24
##	177226	C552049	22082	RIBBON REEL STRIPES DESIGN	-10
##	177227	C552049	22081	RIBBON REEL FLORA + FAUNA	-10
##	177228	C552049	22079	RIBBON REEL HEARTS DESIGN	-10
##		InvoiceDate	UnitPrice	CustomerID	Country
##	236	2010-12-01 09:24:00	1.65	17548	United Kingdom
##	237	2010-12-01 09:24:00	0.29	17548	United Kingdom
##	238	2010-12-01 09:24:00	0.29	17548	United Kingdom
##	239	2010-12-01 09:24:00	0.29	17548	United Kingdom
##	240	2010-12-01 09:24:00	3.45	17548	United Kingdom
##	241	2010-12-01 09:24:00	1.65	17548	United Kingdom
##	242	2010-12-01 09:24:00	1.65	17548	United Kingdom
##	165025	2011-04-20 11:01:00	1.25	17548	United Kingdom
##	165026	2011-04-20 11:01:00	1.65	17548	United Kingdom
##	165027	2011-04-20 11:01:00	1.65	17548	United Kingdom
##	165028	2011-04-20 11:01:00	1.65	17548	United Kingdom
##	165029	2011-04-20 11:01:00	5.95	17548	United Kingdom
##	177224	2011-05-06 09:00:00	5.95	17548	United Kingdom
##	177225	2011-05-06 09:00:00	1.25	17548	United Kingdom
##	177226	2011-05-06 09:00:00	1.65	17548	United Kingdom
##	177227	2011-05-06 09:00:00	1.65	17548	United Kingdom
##	177228	2011-05-06 09:00:00	1.65	17548	United Kingdom

I decide not to delete cancellation invoice rows in general because the influence the quantity and the amount actually spend by each customer. However, I decide to remove two exceptional orders with Quantity larger than 70000 because they exceptionally large, they were clearly mistakes that were cancelled within 15 minutes, and I do not want them to skew the data:

```
onlineRetail = onlineRetail[abs(onlineRetail$Quantity) < 40000, ]
```

## Quantity

All quantities are integer numbers as expected as shown by the following zero count:

```
nrow(onlineRetail[onlineRetail$Quantity != as.integer(onlineRetail$Quantity), ])
```

```
## [1] 0
```

and no quantity is zero:

```
nrow(onlineRetail[onlineRetail$Quantity == 0, ])
```

```
## [1] 0
```

As seen earlier the lowest quantity is negative and the highest quantity is the positive of the same number. Furthermore, the two largest quantities, 80995 and 74215, are cancelled and I decide to leave the data as is.

```
onlineRetail[2500 < abs(onlineRetail$Quantity), ]
```

##	InvoiceNo	StockCode	Description	Quantity
## 4288	C536757	84347	ROTATING SILVER ANGELS T-LIGHT HLDR	-9360
## 4946	536830	84077	WORLD WAR 2 GLIDERS ASSTD DESIGNS	2880
## 52712	540815	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	3114
## 80743	543057	84077	WORLD WAR 2 GLIDERS ASSTD DESIGNS	2592
## 97433	544612	22053	EMPIRE DESIGN ROSETTE	3906
## 160146	C550456	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	-3114
## 160547	550461	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	3114
## 201150	554272	21977	PACK OF 60 PINK PAISLEY CAKE CASES	2700
## 206122	554868	22197	SMALL POPCORN HOLDER	4300
## 270886	560599	18007	ESSENTIAL BALM 3.5g TIN IN ENVELOPE	3186
## 291250	562439	84879	ASSORTED COLOUR BIRD ORNAMENT	2880
## 421633	573008	84077	WORLD WAR 2 GLIDERS ASSTD DESIGNS	4800
## 433789	573995	16014	SMALL CHINESE STYLE SCISSOR	3000
## 502123	578841	84826	ASSTD DESIGN 3D PAPER STICKERS	12540
##	InvoiceDate	UnitPrice	CustomerID	Country
## 4288	2010-12-02 13:23:00	0.03	15838	United Kingdom
## 4946	2010-12-02 16:38:00	0.18	16754	United Kingdom
## 52712	2011-01-11 11:55:00	2.10	15749	United Kingdom
## 80743	2011-02-03 09:50:00	0.21	16333	United Kingdom
## 97433	2011-02-22 09:43:00	0.82	18087	United Kingdom
## 160146	2011-04-18 12:08:00	2.10	15749	United Kingdom
## 160547	2011-04-18 12:20:00	2.10	15749	United Kingdom
## 201150	2011-05-23 12:08:00	0.42	12901	United Kingdom
## 206122	2011-05-27 09:52:00	0.72	13135	United Kingdom
## 270886	2011-07-19 16:04:00	0.06	14609	United Kingdom
## 291250	2011-08-04 17:06:00	1.45	12931	United Kingdom
## 421633	2011-10-27 11:26:00	0.21	12901	United Kingdom
## 433789	2011-11-02 10:24:00	0.32	16308	United Kingdom
## 502123	2011-11-25 14:57:00	0.00	13256	United Kingdom

## InvoiceDate

The following example of the lines making up a full invoice can have different timestamps. It looks like the rows have been keyed-in by hand and that it took a little while for each row:

```
onlineRetail[onlineRetail$InvoiceNo == 544926, ]
```

##	InvoiceNo	StockCode	Description	Quantity
## 101531	544926	37450	CERAMIC CAKE BOWL + HANGING CAKES	6
## 101532	544926	79321	CHILLI LIGHTS	4
## 101533	544926	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6

##	101534	544926	22993	SET OF 4 PANTRY JELLY MOULDS	12
##	101535	544926	48184	DOORMAT ENGLISH ROSE	2
##	101536	544926	48185	DOORMAT FAIRY CAKE	2
##		InvoiceDate	UnitPrice	CustomerID	Country
##	101531	2011-02-24 16:50:00	2.95	13468	United Kingdom
##	101532	2011-02-24 16:50:00	5.75	13468	United Kingdom
##	101533	2011-02-24 16:50:00	2.95	13468	United Kingdom
##	101534	2011-02-24 16:50:00	1.25	13468	United Kingdom
##	101535	2011-02-24 16:50:00	7.95	13468	United Kingdom
##	101536	2011-02-24 16:51:00	7.95	13468	United Kingdom

The date, however, is constant on all of the invoices, as is seen from the following sorted list of the invoices having the most different dates, and they show that all invoices have at most 1 date:

```
onlineRetail %>%
  group_by(InvoiceNo) %>%
  summarize(numberOfDates = length(unique(as.Date(InvoiceDate, "%Y-%m-%d")))) %>%
  arrange(desc(numberOfDates)) %>%
  head(5)
```

```
## # A tibble: 5 x 2
##   InvoiceNo numberOfDates
##   <chr>           <int>
## 1 536365             1
## 2 536366             1
## 3 536367             1
## 4 536368             1
## 5 536369             1
```

## UnitPrice

No unit prices are negative but 33 are zero:

```
onlineRetail[onlineRetail$UnitPrice <= 0, ]
```

##	InvoiceNo	StockCode	Description	Quantity
##	9303	537197 22841	ROUND CAKE TIN VINTAGE GREEN	1
##	33577	539263 22580	ADVENT CALENDAR GINGHAM SACK	4
##	40090	539722 22423	REGENCY CAKESTAND 3 TIER	10
##	47069	540372 22090	PAPER BUNTING RETROSPOT	24
##	47071	540372 22553	PLASTERS IN TIN SKULLS	24
##	56675	541109 22168	ORGANISER WOOD ANTIQUE WHITE	1
##	86790	543599 84535B	FAIRY CAKES NOTEBOOK A6 SIZE	16
##	130189	547417 22062	CERAMIC BOWL WITH LOVE HEART DESIGN	36
##	139454	548318 22055	MINI CAKE STAND HANGING STRAWBERRY	5
##	145209	548871 22162	HEART GARLAND RUSTIC PADDED	2
##	157043	550188 22636	CHILDS BREAKFAST SET CIRCUS PARADE	1
##	187614	553000 47566	PARTY BUNTING	4
##	198384	554037 22619	SET OF 6 SOLDIER SKITTLES	80
##	279325	561284 22167	OVAL WALL MIRROR DIAMANTE	1
##	282913	561669 22960	JAM MAKING SET WITH JARS	11
##	298055	562973 23157	SET OF 6 NATIVITY MAGNETS	240

##	314746	564651	23270	SET OF 2 CERAMIC PAINTED HEARTS	96
##	314747	564651	23268	SET OF 2 CERAMIC CHRISTMAS REINDEER	192
##	314748	564651	22955	36 FOIL STAR CAKE CASES	144
##	314749	564651	21786	POLKADOT RAIN HAT	144
##	379914	569716	22778	GLASS CLOCHE SMALL	2
##	420405	572893	21208	PASTEL COLOUR HONEYCOMB FAN	5
##	436429	574138	23234	BISCUIT TIN VINTAGE CHRISTMAS	216
##	436598	574175	22065	CHRISTMAS PUDDING TRINKET POT	12
##	439362	574469	22385	JUMBO BAG SPACEBOY DESIGN	12
##	446126	574879	22625	RED KITCHEN SCALES	2
##	446794	574920	22899	CHILDREN'S APRON DOLLY GIRL	1
##	446795	574920	23480	MINI LIGHTS WOODLAND MUSHROOMS	1
##	454464	575579	22437	SET OF 9 BLACK SKULL BALLOONS	20
##	454465	575579	22089	PAPER BUNTING VINTAGE PAISLEY	24
##	479080	577129	22464	HANGING METAL HEART LANTERN	4
##	480650	577314	23407	SET OF 2 TRAYS HOME SWEET HOME	2
##	502123	578841	84826	ASSTD DESIGN 3D PAPER STICKERS	12540
##		InvoiceDate	UnitPrice	CustomerID	Country
##	9303	2010-12-05 13:02:00	0	12647	Germany
##	33577	2010-12-16 13:36:00	0	16560	United Kingdom
##	40090	2010-12-21 12:45:00	0	14911	EIRE
##	47069	2011-01-06 15:41:00	0	13081	United Kingdom
##	47071	2011-01-06 15:41:00	0	13081	United Kingdom
##	56675	2011-01-13 14:10:00	0	15107	United Kingdom
##	86790	2011-02-10 12:08:00	0	17560	United Kingdom
##	130189	2011-03-23 09:25:00	0	13239	United Kingdom
##	139454	2011-03-30 11:45:00	0	13113	United Kingdom
##	145209	2011-04-04 13:42:00	0	14410	United Kingdom
##	157043	2011-04-14 17:57:00	0	12457	Switzerland
##	187614	2011-05-12 14:21:00	0	17667	United Kingdom
##	198384	2011-05-20 13:13:00	0	12415	Australia
##	279325	2011-07-26 11:24:00	0	16818	United Kingdom
##	282913	2011-07-28 16:09:00	0	12507	Spain
##	298055	2011-08-11 10:42:00	0	14911	EIRE
##	314746	2011-08-26 13:19:00	0	14646	Netherlands
##	314747	2011-08-26 13:19:00	0	14646	Netherlands
##	314748	2011-08-26 13:19:00	0	14646	Netherlands
##	314749	2011-08-26 13:19:00	0	14646	Netherlands
##	379914	2011-10-06 07:17:00	0	15804	United Kingdom
##	420405	2011-10-26 13:36:00	0	18059	United Kingdom
##	436429	2011-11-03 10:26:00	0	12415	Australia
##	436598	2011-11-03 10:47:00	0	14110	United Kingdom
##	439362	2011-11-04 10:55:00	0	12431	Australia
##	446126	2011-11-07 12:22:00	0	13014	United Kingdom
##	446794	2011-11-07 15:34:00	0	13985	United Kingdom
##	446795	2011-11-07 15:34:00	0	13985	United Kingdom
##	454464	2011-11-10 10:49:00	0	13081	United Kingdom
##	454465	2011-11-10 10:49:00	0	13081	United Kingdom
##	479080	2011-11-17 18:52:00	0	15602	United Kingdom
##	480650	2011-11-18 12:23:00	0	12444	Norway
##	502123	2011-11-25 14:57:00	0	13256	United Kingdom

Apparently the company gives away free stuff from time to time. If the items were actually ordered by the customer and then given free to the customer I would want them to stay in the data to reflect the

customers intention. If the items were not ordered by the customer but just given as a free promotional gift they say nothing about the customer and I would delete them from the data. The problem is that I do not know which one is true. Because there are only 33 invoice rows where the unit price is zero I decide to leave them in the data.

## CustomerID

The data contains 4362 different customer ID's. The customer ID's are all five digit integer as said in the documentation of the data and there is no need to clean-up anything here:

```
nrow(onlineRetail[onlineRetail$CustomerID != as.integer(onlineRetail$CustomerID), ])
```

```
## [1] 0
```

```
nrow(onlineRetail[onlineRetail$CustomerID < 10000 | 99999 < onlineRetail$CustomerID, ])
```

```
## [1] 0
```

Each invoice row has a customer and it could be that invoice rows making up a full invoice would have several different customers. This is not so, however, and it is easy to see, that all invoices have no more than one customer. The following is a sorted list of the invoices having the most different customers, and they show that all invoices have at most 1 customer:

```
onlineRetail %>%
  group_by(InvoiceNo) %>%
  summarize(numberOfCustomers = length(unique(CustomerID))) %>%
  arrange(desc(numberOfCustomers)) %>%
  head(5)
```

```
## # A tibble: 5 x 2
##   InvoiceNo numberOfCustomers
##   <chr>          <int>
## 1 536365          1
## 2 536366          1
## 3 536367          1
## 4 536368          1
## 5 536369          1
```

## Country

The countries include an 'Unspecified' and all seems Ok except that it also includes 'European Community' that includes several of the other countries:

```
unique(onlineRetail$Country)
```

```
## [1] "United Kingdom"      "France"              "Australia"
## [4] "Netherlands"         "Germany"             "Norway"
## [7] "EIRE"                "Switzerland"         "Spain"
## [10] "Poland"              "Portugal"            "Italy"
## [13] "Belgium"             "Lithuania"           "Japan"
```

```
## [16] "Iceland"           "Channel Islands"    "Denmark"
## [19] "Cyprus"            "Sweden"             "Austria"
## [22] "Israel"           "Finland"            "Greece"
## [25] "Singapore"        "Lebanon"            "United Arab Emirates"
## [28] "Saudi Arabia"     "Czech Republic"     "Canada"
## [31] "Unspecified"      "Brazil"             "USA"
## [34] "European Community" "Bahrain"            "Malta"
## [37] "RSA"
```

It it seen, though, that only one customer (with four invoices) has ‘European Community’ as country and I decide to leave it in the data:

```
onlineRetail[onlineRetail$Country == 'European Community', ] %>%
  distinct(InvoiceNo, CustomerID)
```

```
##   InvoiceNo CustomerID
## 1    551013      15108
## 2    555542      15108
## 3    C556294      15108
## 4     560783      15108
```

## The Data Types

Finally, I change some of the fields from being strings or numerics to be categorial data. For example the `CustomerID` that is read as a number from the original data, but it of course make no sense to do arithmetic on customer ID’s, you cannot add two customer ID’s for example. And all in all I end up with the following data ready to be analysed:

```
## Observations: 404,905
## Variables: 8
## $ InvoiceNo    <fct> 536365, 536365, 536365, 536365, 536365, 536365, 53...
## $ StockCode   <fct> 85123A, 71053, 84406B, 84029G, 84029E, 22752, 2173...
## $ Description <chr> "WHITE HANGING HEART T-LIGHT HOLDER", "WHITE METAL...
## $ Quantity    <int> 6, 6, 8, 6, 6, 2, 6, 6, 6, 32, 6, 6, 8, 6, 6, 3, 2...
## $ InvoiceDate  <dtm> 2010-12-01 07:26:00, 2010-12-01 07:26:00, 2010-12...
## $ UnitPrice   <dbl> 2.55, 3.39, 2.75, 3.39, 3.39, 7.65, 4.25, 1.85, 1....
## $ CustomerID  <fct> 17850, 17850, 17850, 17850, 17850, 17850, 17850, 1...
## $ Country     <fct> UnitedKingdom, UnitedKingdom, UnitedKingdom, Unite...
```

## Sizes after Cleaning the Data

After cleaning the dataset it contains 21784 different invoices and 4362 different customers.

# Part 1: Analysis of Customer RFM and Customer Segments

## Introduction

RFM is an often used way to measure the value of a customer. The three letters stand for **R**ecency, **F**requency, and **M**onetary Value. For the dataset `onlineRetail` I use the following:

- Recency: the number of days before the last date of the period
- Frequency: the number of different invoices and cancellations over the period
- Monetary Value: the total amount paid over the period (does take cancellations inside the period into account)

That is, I take the full year-long dataset as the period to compute for. In a real business setting you would get new data every day, and you could periodically recalculate the results I get in this report on data for the last year or two.

My goal is to cluster the customers from their RFM, first using a traditional, fixed algorithm, and second using a machine learning clustering algorithm and finally compare the two customer segmentations.

## Calculate RFM for each Customer

I build a new dataset `CustomerRFM` that contains the RFM values for each customer from the `onlineRetail` dataset:

```
RecencyDate = as.Date(max(onlineRetail$InvoiceDate))
CustomerRFM <- onlineRetail %>%
  group_by(CustomerID) %>%
  summarize(Recency = as.numeric(RecencyDate - as.Date(max(InvoiceDate))),
            Frequency = n_distinct(InvoiceNo),
            Monetary = sum(Quantity * UnitPrice))
head(CustomerRFM)
```

```
## # A tibble: 6 x 4
##   CustomerID Recency Frequency Monetary
##   <fct>      <dbl>      <int>      <dbl>
## 1 12347         2         7        4310
## 2 12348        75         4       1437.
## 3 12349        18         1       1458.
## 4 12350       310         1        294.
## 5 12352        36         8       1265.
## 6 12353       204         1         89
```

```
summary(CustomerRFM)
```

```
##   CustomerID      Recency      Frequency      Monetary
## 12347 : 1   Min.   : 0.00   Min.   : 1.000   Min.   : -1192.2
## 12348 : 1   1st Qu.: 16.00   1st Qu.: 1.000   1st Qu.:  294.3
## 12349 : 1   Median : 50.00   Median : 3.000   Median :   645.7
## 12350 : 1   Mean    : 91.68   Mean    : 4.994   Mean    :  1899.7
## 12352 : 1   3rd Qu.:143.00   3rd Qu.: 5.000   3rd Qu.:  1596.4
## 12353 : 1   Max.    :373.00   Max.    :242.000   Max.    :278778.0
## (Other):4356
```

I need to scale the values before using them. If for example the first feature in a dataset ranges from 0 to 1 and the second feature ranges from 0 to 1000, the computed Euclidean distance between two points, or rows, from the dataset will be nearly totally dominated by the second feature. To avoid this problem the features must be scaled to the same ranges before using them in a clustering algorithm.

I use the following to scale the three feature R, F, and M. The method I use assumes that the values in each feature is close to be normal distributed and scale the values to have mean 0 and standard deviation 1. The values in `CustomerRFM` are not normally distributed but scaling this way is much better than not scaling:



```
CustomerRFM$Recency <- scale(CustomerRFM$Recency)
CustomerRFM$Frequency <- scale(CustomerRFM$Frequency)
CustomerRFM$Monetary <- scale(CustomerRFM$Monetary)
head(CustomerRFM)
```

```
## # A tibble: 6 x 4
##   CustomerID Recency[,1] Frequency[,1] Monetary[,1]
##   <fct>          <dbl>          <dbl>          <dbl>
## 1 12347          -0.889           0.221           0.292
## 2 12348          -0.165          -0.109          -0.0560
## 3 12349          -0.730          -0.440          -0.0535
## 4 12350           2.16          -0.440          -0.194
## 5 12352          -0.552           0.331          -0.0767
## 6 12353           1.11          -0.440          -0.219
```

And we end as expected up with scaled values for each feature:

```
## [1] "Recency: mean = -0.000000, sd = 1.000000"

## [1] "Frequency: mean = 0.000000, sd = 1.000000"

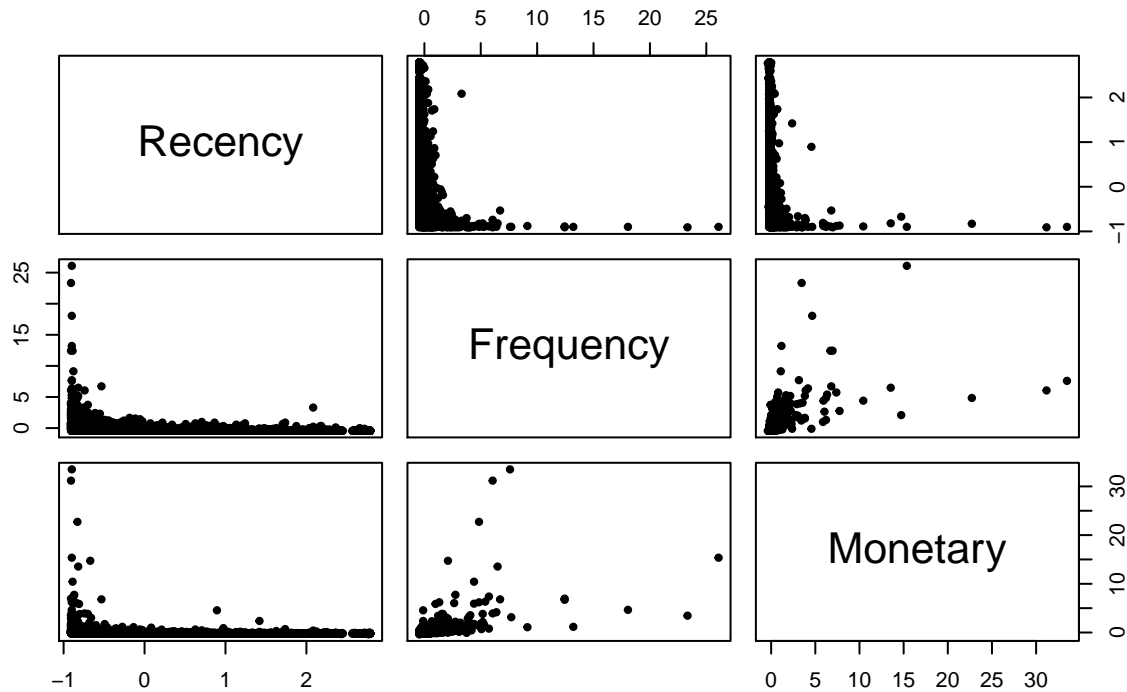
## [1] "Monetary: mean = 0.000000, sd = 1.000000"
```

## Visually Illustrating the Values and Principal Component Analysis

In the following plot it is seen that all three features are pairwise correlated but not in a simple, for example, linear way. Recency, though, seems to be correlated in somewhat the same way to both Frequency and Monetary which suggests that much of the variability in the data can be expressed with only two features, if we can find the suitable features.

```
pairs(~Recency+Frequency+Monetary,data=CustomerRFM,
      main="Simple Scatterplot Matrix", pch = 20)
```


## Simple Scatterplot Matrix



To illustrate the RFM values for the customers I first draw a three-dimensional plot where each point is a customer and each of the three coordinates is one of the three RFM features such that the points position is the RFM values of that customer.

```
scatterplot3d(CustomerRFM$Recency, CustomerRFM$Frequency, CustomerRFM$Monetary,  
              pch = 20,  
              main="3D Scatterplot",  
              xlab = "Recency", ylab = "Frequency", zlab = "Monetary")
```

### 3D Scatterplot



It is worth noting that a big part of the space is empty. For example there are no customers that buy frequently but just not recently. A possible explanation is that frequency and recency is connected in that if a customer has not bought in a while then the frequency falls.

It is hard, at least for me, to see the exact three-dimensional positions of the points in this three-dimensional plot. It would be much easier to look at in a turnable, and possibly interactive, three-dimensional plot but this is not possible in the pdf format of this report.

So in addition to the three dimensional plot I also illustrate the positions of the points in a two-dimensional plot where I map all the points from three dimensions into two dimensions using Principal Component Analysis (PCA). PCA computes a new three-dimensional coordinate system where each coordinate, in order, covers as much variability as possible. So dropping the least important coordinate after PCA will usually give a much better result, and never be worse, than just dropping one of the coordinates without PCA.

```
CustomerRFM_PCA <- prcomp(CustomerRFM[c(2,3,4)])
summary(CustomerRFM_PCA)
```

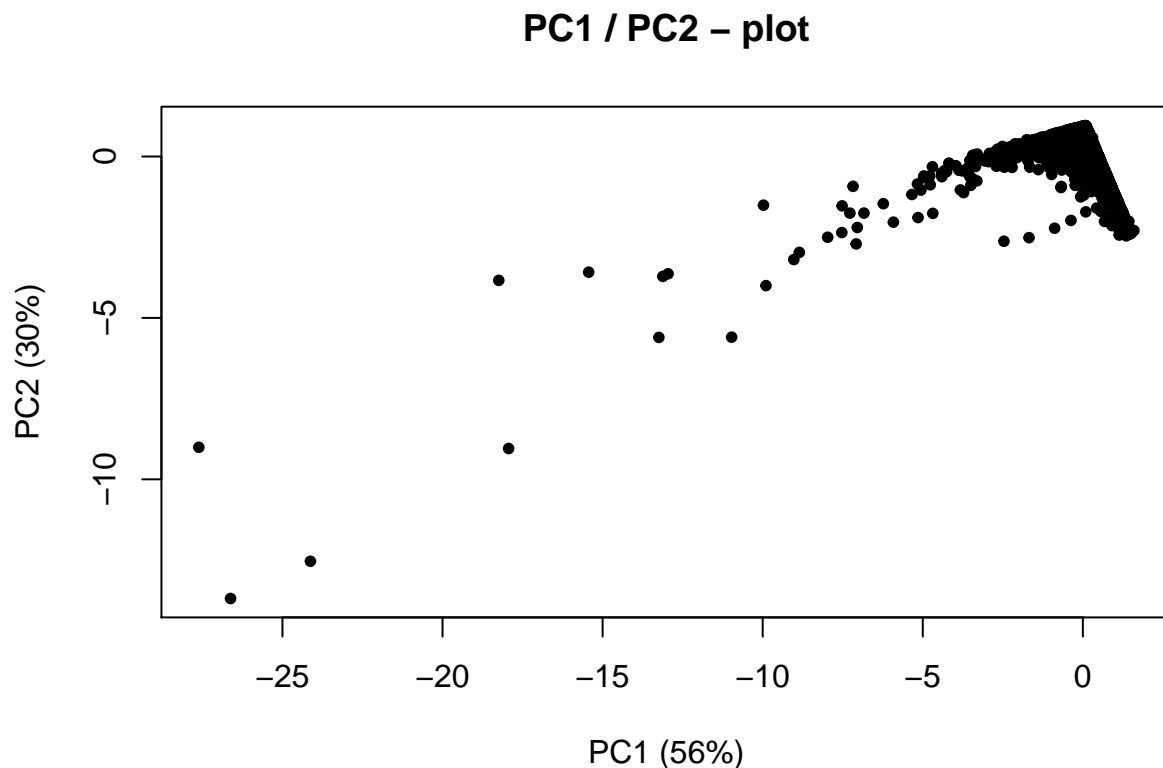
```
## Importance of components:
##              PC1      PC2      PC3
## Standard deviation  1.2928 0.9498 0.6532
## Proportion of Variance 0.5571 0.3007 0.1422
## Cumulative Proportion 0.5571 0.8578 1.0000
```

The three new dimensions are sorted by the amount of variability of the data in the coordinates direction (*standard deviation*). It is seen that the two first coordinates together explain close to 86% of the variance of the data, which means that using the first two coordinates will give a reasonable illustration of the customers

RFM values. However, it is an approximation only and there is no need to reduce the number of features from three to two in the further calculations. So I use the PCA for illustration only.

In the following plot it is easy to see, that most of the customers clusters together in one big angular formed cluster with fewer customers radiating out of the big cluster. I had hoped that a small number, five to ten, clusters formed by the positions of the points, i.e. customers, were clearly visible. This is however not the case and any customer segmentation will only be a rough and simplifying description of the much more continuously distributed customers.

```
plot(CustomerRFM_PCA$x[,1], CustomerRFM_PCA$x[,2],  
      xlab="PC1 (56%)", ylab = "PC2 (30%)",  
      main = "PC1 / PC2 - plot",  
      pch = 20, col = "black")
```



## Customer Segmentation using a Traditional Fixed Method

The Internet contains many sources that explain their way to split customers into Customer Segment using their RFM values. While there seems to be many different ways to segment customers based on RFM, all I have found is based on the same principles and differ only on the details and on the definition of the final customer segments. I use the process from <https://www.putler.com/rfm-analysis/> because their method is close to all other methods that I found and they describe it detailed.

First a score 1 to 5 is assigned to each of R, F, and M using quintiles, i.e. it scores 5 if the value is ranked in the best 20%, it scores 4 if ranked in the best 20-40%, etc. For R lower values are best, for F and M higher values are best. I do not need to scale the original RFM values but because the rank is not changed by scaling it does not matter and I use the scaled `CustomerRFM` that I already have build.

```
n = nrow(CustomerRFM)
CustomerRFM$RecencyScore <- as.integer(1 + 5 * (1 - rank(CustomerRFM$Recency) / n))
CustomerRFM$FrequencyScore <- as.integer(1 + 5 * rank(CustomerRFM$Frequency) / n)
CustomerRFM$MonetaryScore <- as.integer(1 + 5 * rank(CustomerRFM$Monetary) / n)
head(CustomerRFM)
```

```
## # A tibble: 6 x 7
##   CustomerID Recency[,1] Frequency[,1] Monetary[,1] RecencyScore
##   <fct>          <dbl>          <dbl>          <dbl>          <int>
## 1 12347          -0.889           0.221           0.292           5
## 2 12348          -0.165          -0.109          -0.0560          2
## 3 12349          -0.730          -0.440          -0.0535          4
## 4 12350           2.16          -0.440          -0.194           1
## 5 12352          -0.552           0.331          -0.0767          3
## 6 12353           1.11          -0.440          -0.219           1
## # ... with 2 more variables: FrequencyScore <int>, MonetaryScore <int>
```

The customer are now assigned to one of eleven customer segments dependent on their R, F, and M scores. For example customers with R, F, and M scores of 5, 5, and 5 are in the Champions customer segment - they bought recently, buy often and spend a lot. The eleven customer segments are called: “Champions”, “Loyal Customers”, “Potential Loyalist”, “Recent Customers”, “Promising”, “Customers Needing Attention”, “About To Sleep”, “At Risk”, “Can’t Lose Them”, “Hibernating”, “Lost”.

Most of the eleven customer segments combine different R, F, and M scores. The following function computes them all:

```
putlerSegment <- function(row) {
  fm <- as.integer((as.integer(row["FrequencyScore"]) + as.integer(row["MonetaryScore"])) / 2) - 1
  idx = 1 + 5 * (as.integer(row["RecencyScore"]) - 1) + fm
  switch(idx,
    "Lost",
    "Lost",
    "AtRisk",
    "AtRisk",
    "CantLoseThem",
    "Lost",
    "Hibernating",
    "AtRisk",
    "AtRisk",
    "AtRisk",
    "AboutToSleep",
    "AboutToSleep",
    "NeedingAttention",
    "LoyalCustomers",
    "LoyalCustomers",
    "Promising",
    "PotentialLoyalist",
    "PotentialLoyalist",
    "LoyalCustomers",
    "LoyalCustomers",
    "RecentCustomers",
    "PotentialLoyalist",
    "PotentialLoyalist",
```

```

    "LoyalCustomers",
    "Champions")
}

```

and after computing the Putler Customer Segments the data looks like:

```

CustomerRFM$putlerSegment <- apply(CustomerRFM, 1, putlerSegment)
head(CustomerRFM)

## # A tibble: 6 x 8
##   CustomerID Recency[,1] Frequency[,1] Monetary[,1] RecencyScore
##   <fct>          <dbl>          <dbl>          <dbl>          <int>
## 1 12347          -0.889           0.221           0.292           5
## 2 12348          -0.165           -0.109          -0.0560          2
## 3 12349          -0.730           -0.440          -0.0535          4
## 4 12350           2.16           -0.440          -0.194           1
## 5 12352          -0.552           0.331          -0.0767          3
## 6 12353           1.11           -0.440          -0.219           1
## # ... with 3 more variables: FrequencyScore <int>, MonetaryScore <int>,
## #   putlerSegment <chr>

```

Cleaning up the sums and balances:

```

CustomerRFM$RecencyScore <- NULL
CustomerRFM$FrequencyScore <- NULL
CustomerRFM$MonetaryScore <- NULL
CustomerRFM$putlerSegment <- as.factor(CustomerRFM$putlerSegment)

head(CustomerRFM)

## # A tibble: 6 x 5
##   CustomerID Recency[,1] Frequency[,1] Monetary[,1] putlerSegment
##   <fct>          <dbl>          <dbl>          <dbl> <fct>
## 1 12347          -0.889           0.221           0.292 Champions
## 2 12348          -0.165           -0.109          -0.0560 AtRisk
## 3 12349          -0.730           -0.440          -0.0535 PotentialLoyalist
## 4 12350           2.16           -0.440          -0.194 Lost
## 5 12352          -0.552           0.331          -0.0767 LoyalCustomers
## 6 12353           1.11           -0.440          -0.219 Lost

```

The following two plots illustrate the customers Putler segments with each segment being a different color:

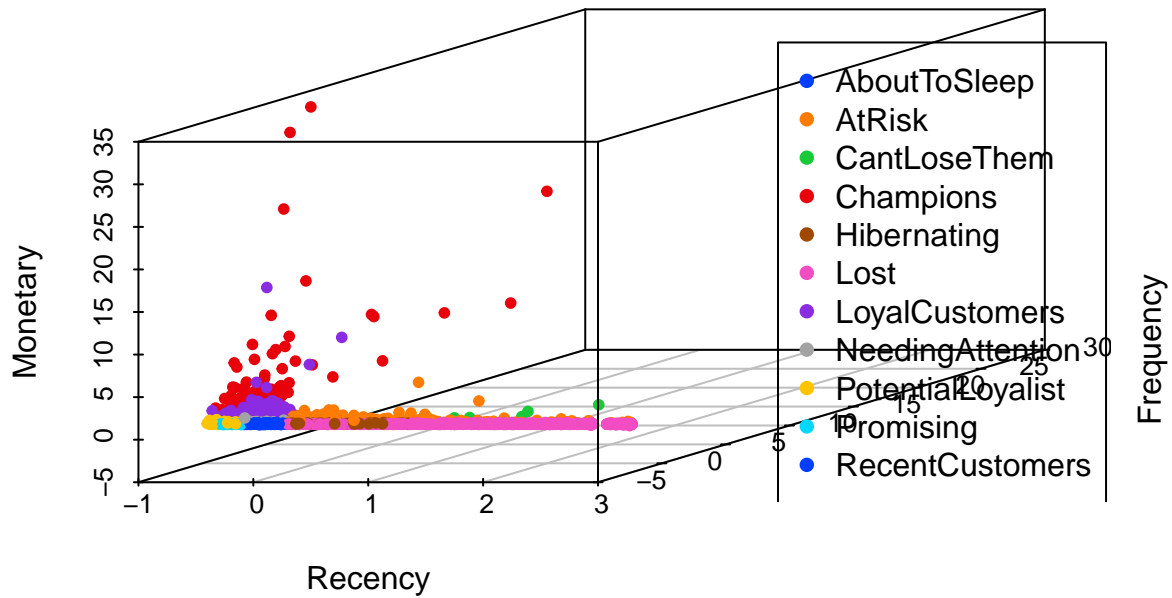
```

colors <- colorpalette[CustomerRFM$putlerSegment]

plot3d <- scatterplot3d(CustomerRFM$Recency, CustomerRFM$Frequency, CustomerRFM$Monetary,
  xlab = "Recency", ylab = "Frequency", zlab = "Monetary",
  main="Putler Customer Segments", color = colors, pch = 20)
legend(plot3d$xyz.convert(2.9, 10, 40),
  legend = levels(CustomerRFM$putlerSegment),
  col = colorpalette, pch = 16)

```

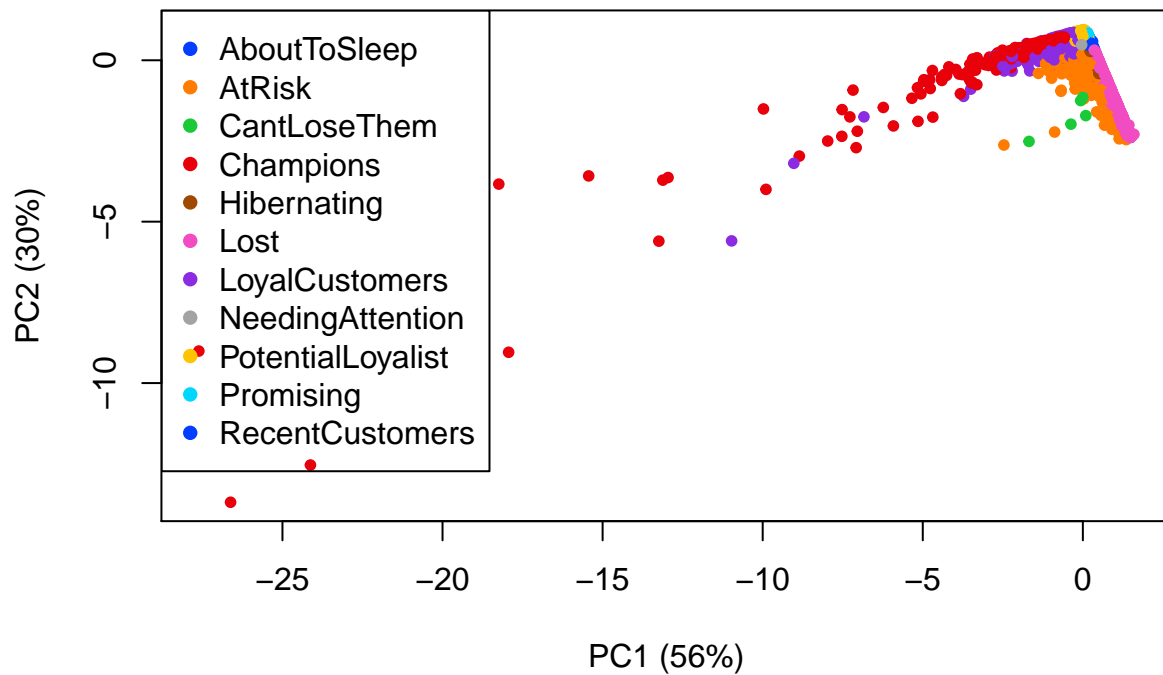
## Putler Customer Segments



```
colors <- colorpalette[CustomerRFM$putlerSegment]

plot(CustomerRFM_PCA$x[,1], CustomerRFM_PCA$x[,2],
      xlab="PC1 (56%)", ylab = "PC2 (30%)",
      main = "Putler Customer Segments (mapped to 2D)",
      pch = 20, col = colors)
legend("topleft", legend = levels(CustomerRFM$putlerSegment),
      col = colorpalette, pch = 16)
```

## Putler Customer Segments (mapped to 2D)



The plots show that the Putler segments seems to differ a lot in how many customers are in each segment and especially the Champions segment seems to cover a much larger area than the other segments.

The number of customers in each Putler segment is:

```
CustomerRFM %>%
  group_by(putlerSegment) %>%
  summarize(NumberOfCustomers = n()) %>%
  arrange(desc(NumberOfCustomers))
```

```
## # A tibble: 11 x 2
##   putlerSegment   NumberOfCustomers
##   <fct>           <int>
## 1 Lost            1020
## 2 LoyalCustomers    898
## 3 PotentialLoyalist 548
## 4 AtRisk           465
## 5 AboutToSleep     435
## 6 Champions        354
## 7 Hibernating      254
## 8 NeedingAttention  162
## 9 Promising        155
## 10 RecentCustomers   66
## 11 CantLoseThem      5
```



## Customer Segmentation using K-means Machine Learning Clustering Method


Clustering is an unsupervised machine learning technique which means that no examples of a correct solution are given beforehand. Instead the goal is to find clusters in the given data using some measure of distance. Several different algorithms exist to find clusters. I choose to use the K-means algorithm with Euclidean distance because it is easy to understand and fast enough for the use here.

K-means starts by selecting k cluster centers randomly or in some predefined pattern. In step 1 all the data points are assigned to its closest cluster center. In step 2 the cluster centers are moved to the mean of all the data points in the cluster. Step1 and 2 are repeated until the cluster centers have converged or a maximal predefined number of steps have been executed. Actually there are implementations of several variations of the basic algorithm but I just use the default.

First I want to find the optimal number of clusters to use. The goal is to have high similarity within each cluster and low similarity across clusters and I use the so-called elbow method where I plot the percentage of explained variance against the number of clusters. At a certain point adding one more cluster does not increase the percentage of explained variance as much as before. This is the so-called elbow point and the number of clusters at the elbow point is optimal in this sense.

```
set.seed(9)
pExplainedVar = rep(0,20)
df <- select(CustomerRFM, Recency, Frequency, Monetary)
for (i in 2:20) {
  fit = kmeans(df, centers = i, iter.max = 15, nstart = 20)
  pExplainedVar[i] = 100 * fit$betweenss / fit$totss
}

plot(1:20, pExplainedVar, type="b", pch = 20, xlab="Number of Clusters", ylab="Percentage of Explained Variance")
text(1:20, pExplainedVar, labels = 1:20, cex = 0.7, pos = 3)
```



From the above plot I think that the flattening-out of the increase begins somewhere around seven or eight and I choose to continue to find seven clusters using the K-means algorithm.

```
set.seed(9)
df <- select(CustomerRFM, Recency, Frequency, Monetary)
k_result <- kmeans(df, centers = 7, iter.max = 15, nstart = 25)

k_result$centers
```

```
##      Recency  Frequency  Monetary
## 1 -0.7583631  1.3418178  0.41970033
## 2 -0.5613750 -0.1462504 -0.09410573
## 3 -0.8054969  4.0383745  4.18725744
## 4 -0.8787235  6.1646198 29.13352530
## 5  2.0342762 -0.3861998 -0.18589284
## 6  0.6569728 -0.2858412 -0.14727425
## 7 -0.9001920 17.5936359  6.38494369
```

Adding the found clusters to the CustomerRFM and draw the same plots as for the Putler segments:

```
CustomerRFM$kmeansSegment <- as.factor(k_result$cluster)
head(CustomerRFM)
```

```
## # A tibble: 6 x 6
##   CustomerID Recency[,1] Frequency[,1] Monetary[,1] putlerSegment
```


```
##      <fct>          <dbl>      <dbl>      <dbl> <fct>
## 1 12347          -0.889       0.221      0.292 Champions
## 2 12348          -0.165      -0.109     -0.0560 AtRisk
## 3 12349          -0.730      -0.440     -0.0535 PotentialLoy~
## 4 12350           2.16       -0.440     -0.194  Lost
## 5 12352          -0.552       0.331     -0.0767 LoyalCustome~
## 6 12353           1.11       -0.440     -0.219  Lost
## # ... with 1 more variable: kmeansSegment <fct>
```

The following two plots illustrate the customers K-means segments with each segment being a different color:

```
colors <- colorpalette[CustomerRFM$kmeansSegment]

plot3d <- scatterplot3d(CustomerRFM$Recency, CustomerRFM$Frequency, CustomerRFM$Monetary,
                        xlab = "Recency", ylab = "Frequency", zlab = "Monetary",
                        main="K-means Customer Segments", color = colors, pch = 20)
legend(plot3d$xyz.convert(2.9, 10, 40),
       legend = levels(CustomerRFM$kmeansSegment),
       col = colorpalette, pch = 20)
```


## K-means Customer Segments



```
colors <- colorpalette[CustomerRFM$kmeansSegment]

plot(CustomerRFM_PCA$x[,1], CustomerRFM_PCA$x[,2],
     xlab="PC1 (56%)", ylab = "PC2 (30%)",
     main = "K-means Customer Segments (mapped to 2D)",
```

```
pch = 20, col = colors)
legend("topleft", legend = levels(CustomerRFM$kmeansSegment),
      col = colorpalette, pch = 16)
```



The plots show that the K-means segments seems like the Putler segments to differ a lot in how many customers are in each segment and in how large area they cover.

The number of customers in each K-means segment is:

```
CustomerRFM %>%
  group_by(kmeansSegment) %>%
  summarize(NumberOfCustomers = n()) %>%
  arrange(desc(NumberOfCustomers))
```

```
## # A tibble: 7 x 2
##   kmeansSegment NumberOfCustomers
##   <fct>          <int>
## 1 2              2522
## 2 6              773
## 3 5              616
## 4 1              401
## 5 3              41
## 6 7              6
## 7 4              3
```


To be useful we need a business interpretation of the seven K-means clusters. I make that by looking at three projections of the three dimensional plot of the K-means clusters by leaving out one of the coordinates.

```

colors <- colorpalette[CustomerRFM$kmeansSegment]
plot(CustomerRFM$Recency, CustomerRFM$Frequency,
      xlab = "Recency", ylab = "Frequency",
      main="K-means Customer Segments", col = colors, pch = 20)
legend("topright", legend = levels(CustomerRFM$kmeansSegment),
      col = colorpalette, pch = 16)

```

## K-means Customer Segments




```


colors <- colorpalette[CustomerRFM$kmeansSegment]
plot(CustomerRFM$Recency, CustomerRFM$Monetary,
      xlab = "Recency", ylab = "Monetary",
      main="K-means Customer Segments", col = colors, pch = 20)
legend("topright", legend = levels(CustomerRFM$kmeansSegment),
      col = colorpalette, pch = 16)

```

## K-means Customer Segments



```
colors <- colorpalette[CustomerRFM$kmeansSegment]
plot(CustomerRFM$Frequency, CustomerRFM$Monetary,
      xlab = "Frequency", ylab = "Monetary",
      main="K-means Customer Segments", col = colors, pch = 20)
legend("topright", legend = levels(CustomerRFM$kmeansSegment),
      col = colorpalette, pch = 16)
```



From looking at the three two-dimsional projections of the K-means clusters I suggest the following description of the K-means segments:

Number	Name	Description
1	Small and Occasional	Recently active but spends low and not very often
2	New	Fairly recently active but spendend low and have not been active earlier
3	Good Average	Active and buys from seldom to often and spends from low to average
4	Big Spenders	Spend most money of all and they recently did it again
5	Inactive	Not active for a long time and spendend low
6	Fading Away	Not active for some time and spendend low
7	Good and Loyal	Spend average or more money and buy very often

To compare with the Putler segments I computer for each K-means cluster the percentual composition of customers that are in each of the Putler clusters:

```
CustomerRFM %>%
  group_by(kmeansSegment) %>%
  count(putlerSegment) %>%
  mutate(percent = 100*n/sum(n))
```

```
## # A tibble: 26 x 4
## # Groups:   kmeansSegment [7]
```

```
##      kmeansSegment putlerSegment      n percent
##      <fct>          <fct>          <int>  <dbl>
##    1 1              AtRisk           9    2.24
##    2 1              CantLoseThem      1    0.249
##    3 1              Champions        229   57.1
##    4 1              LoyalCustomers    162   40.4
##    5 2              AboutToSleep      435   17.2
##    6 2              AtRisk           171    6.78
##    7 2              Champions         81    3.21
##    8 2              Hibernating        92    3.65
##    9 2              Lost              81    3.21
##   10 2              LoyalCustomers    731   29.0
## # ... with 16 more rows
```

It is seen from the table that especially K-means segment 7 overlap with many Putler segments, and it is seen that for example the Putler Champions segment is split into several of the K-means segments. All in all the two sets of segments parts the customers in very different ways.

## Part 2: Analysis of Predicting Customer Segments from their First Order

### Introduction

I want to predict in which segment a new customer ends up at the end of the period of the given data. I will use all the given features of the customer available at the time of the the customers first purchase. This a classification problem and it is grouped as a supervised machine learning problem because we train a model from a training set where the correct answer is given.

I train two different models for each of the K-means segment prediction and the Putler segment prediction. First I use k-nearest neighbors and then I use random forest because I want to try more than one algorithm to see how they compare. Also I compare to a simple random approach of guessing the segment randomly according to the segment sizes.

For this investigation I assume that each customers first non-cancellation order is the customers first order, even though I cannot know if the customer has been a customer for years before the first date in my dataset.

### The Class Imbalance Problem

We saw earlier that the number of customers in the Putler segments varies from a few to about 1000, and that the numbers of customers in the Kmeans segments varies from a few to about 2500. The large difference of the class sizes, the imbalanced classes, makes it very possible that the classification algorithms optimises the accuracy at the expense of the precision, and we risk an algorithm that always predicts the most populated category.

Different approaches to prevent this exists, but it is not the purpose of this report to compare different possible solution to the class imbalance problem and I just choose a method that often works and after the training of the models I check if the predictions are spread on several categories. The approach I choose is oversampling of the small categories. It is chosen in the caret train method by adding `sampling = "up"` to the `trainControl`.



## Feature extraction

First I construct a data frame that holds information about the first purchase for each customer. I ignore all the cancellations because I only want the first non-cancellation order for each customer. In the process we loose about twenty customers who only had cancellations in the dataset.

```
firstOrder <- onlineRetail %>%
  filter(!startsWith(tolower(as.character(InvoiceNo)), "c")) %>%
  group_by(CustomerID, InvoiceNo) %>%
  summarize(N = n(), TotalQuantity = sum(Quantity),
            TimeOfDay = (60 * hour(min(InvoiceDate)) + minute(min(InvoiceDate))),
            TotalPrice = sum(Quantity*UnitPrice), Country = first(Country),
            InvoiceDate = min(InvoiceDate)) %>%
  slice(which.min(InvoiceDate)) %>%
  ungroup() %>%
  inner_join(CustomerRFM, by = "CustomerID") %>%
  select(-CustomerID, -Recency, -Frequency, -Monetary, -InvoiceDate, -InvoiceNo)

print(nrow(firstOrder))
```

```
## [1] 4334
```

```
head(firstOrder)
```

```
## # A tibble: 6 x 7
##       N TotalQuantity TimeOfDay TotalPrice Country putlerSegment
##   <int>      <int>      <dbl>      <dbl> <fct>    <fct>
## 1    31         319        837        712. Iceland Champions
## 2    16         1248       1089        653. Finland AtRisk
## 3    72          630        531       1458. Italy   PotentialLoy~
## 4    16          196        901        294. Norway   Lost
## 5    15           98        693        296. Norway   LoyalCustome~
## 6     4           20       1007         89 Bahrain Lost
## # ... with 1 more variable: kmeansSegment <fct>
```

Because I want to use the k-nearest neighbor algorithm, see more about this in a moment, and because the k-nearest neighbor algorithm use the distance between points I need to scale the numeric data before I train the models. Also I change the Country feature to dummy variable, aka one hot encoding, to make it numeric:

```
dummy <- dummyVars("~ Country", data = firstOrder, sep = NULL)
firstOrderEncoded <- cbind(firstOrder, predict(dummy, newdata = firstOrder))

firstOrderEncoded <- firstOrderEncoded %>%
  select(-Country) %>%
  mutate_if(is.numeric, function(c1m) { as.vector(scale(c1m)) })

glimpse(firstOrderEncoded)
```

```
## Observations: 4,334
## Variables: 43
## $ N
      <dbl> 0.29595976, -0.33561666, 2.02226866,...
```

```
## $ TotalQuantity      <dbl> 0.17380512, 2.42427831, 0.92719281, ...
## $ TimeOfDay          <dbl> 0.80532341, 2.59495502, -1.36780069,...
## $ TotalPrice         <dbl> 0.560181157, 0.453026216, 1.91484921...
## $ putlerSegment      <fct> Champions, AtRisk, PotentialLoyalist...
## $ kmeansSegment      <fct> 2, 2, 2, 5, 2, 6, 5, 6, 2, 2, 2, ...
## $ CountryAustralia   <dbl> -0.04561189, -0.04561189, -0.0456118...
## $ CountryAustria     <dbl> -0.04561189, -0.04561189, -0.0456118...
## $ CountryBahrain     <dbl> -0.02148427, -0.02148427, -0.0214842...
## $ CountryBelgium     <dbl> -0.07461341, -0.07461341, -0.0746134...
## $ CountryBrazil      <dbl> -0.01518992, -0.01518992, -0.0151899...
## $ CountryCanada      <dbl> -0.03039037, -0.03039037, -0.0303903...
## $ CountryChannelIslands <dbl> -0.04561189, -0.04561189, -0.0456118...
## $ CountryCyprus       <dbl> -0.04021661, -0.04021661, -0.0402166...
## $ CountryCzechRepublic <dbl> -0.01518992, -0.01518992, -0.0151899...
## $ CountryDenmark     <dbl> -0.04299833, -0.04299833, -0.0429983...
## $ CountryEIRE        <dbl> -0.02631579, -0.02631579, -0.0263157...
## $ CountryEuropeanCommunity <dbl> -0.01518992, -0.01518992, -0.0151899...
## $ CountryFinland     <dbl> -0.05268635, 18.97586794, -0.0526863...
## $ CountryFrance      <dbl> -0.1431095, -0.1431095, -0.1431095, ...
## $ CountryGermany     <dbl> -0.1488781, -0.1488781, -0.1488781, ...
## $ CountryGreece      <dbl> -0.03039037, -0.03039037, -0.0303903...
## $ CountryIceland     <dbl> 65.81793244, -0.01518992, -0.0151899...
## $ CountryIsrael      <dbl> -0.02631579, -0.02631579, -0.0263157...
## $ CountryItaly       <dbl> -0.05692094, -0.05692094, 17.5641746...
## $ CountryJapan       <dbl> -0.04299833, -0.04299833, -0.0429983...
## $ CountryLebanon     <dbl> -0.01518992, -0.01518992, -0.0151899...
## $ CountryLithuania   <dbl> -0.01518992, -0.01518992, -0.0151899...
## $ CountryMalta       <dbl> -0.02148427, -0.02148427, -0.0214842...
## $ CountryNetherlands <dbl> -0.04561189, -0.04561189, -0.0456118...
## $ CountryNorway      <dbl> -0.04808472, -0.04808472, -0.0480847...
## $ CountryPoland      <dbl> -0.03722904, -0.03722904, -0.0372290...
## $ CountryPortugal    <dbl> -0.06634929, -0.06634929, -0.0663492...
## $ CountryRSA         <dbl> -0.01518992, -0.01518992, -0.0151899...
## $ CountrySaudiArabia <dbl> -0.01518992, -0.01518992, -0.0151899...
## $ CountrySingapore   <dbl> -0.01518992, -0.01518992, -0.0151899...
## $ CountrySpain       <dbl> -0.08062912, -0.08062912, -0.0806291...
## $ CountrySweden      <dbl> -0.04299833, -0.04299833, -0.0429983...
## $ CountrySwitzerland <dbl> -0.06808083, -0.06808083, -0.0680808...
## $ CountryUnitedArabEmirates <dbl> -0.02148427, -0.02148427, -0.0214842...
## $ CountryUnitedKingdom <dbl> -3.060435, -3.060435, -3.060435, -3....
## $ CountryUnspecified <dbl> -0.03039037, -0.03039037, -0.0303903...
## $ CountryUSA         <dbl> -0.03039037, -0.03039037, -0.0303903...
```

Using the R caret package I first split my data into test and training sets with 70% of the original data being the training set and 30% being the test set:

```
set.seed(9)
test_index <- createDataPartition(firstOrderEncoded$kmeansSegment, times = 1, p = 0.3, list = FALSE)

test_set <- firstOrderEncoded[test_index, ]
train_set <- firstOrderEncoded[-test_index, ]

cat("Size of test_set =", nrow(test_set), "\n")
```

```
## Size of test_set = 1303
```

```
cat("Size of train_set =", nrow(train_set), "\n")
```

```
## Size of train_set = 3031
```

## Random Guessing for Putler Segments

As a baseline I use a simple predictor that for any customer always predicts randomly according to the sizes of each segment. First for the Putler segments:

```
putlerSegments <- CustomerRFM %>%  
  group_by(putlerSegment) %>%  
  summarize(NumberOfCustomers = n())
```

and when predicting the testset we get the following accuracy:

```
set.seed(9)  
idx <- sample(nrow(putlerSegments), size = nrow(test_set), replace=TRUE, prob = putlerSegments$NumberOfCustomers)  
test_pred <- putlerSegments[idx, ]$putlerSegment  
accuracy <- mean(test_pred == test_set$putlerSegment)  
accuracy_random_onPutler = sprintf("%.2f", accuracy)  
accuracy_random_onPutler
```

```
## [1] "0.13"
```

## Random Guessing for Kmeans Segments

The baseline for the Kmeans segments is:

```
kmeansSegments <- CustomerRFM %>%  
  group_by(kmeansSegment) %>%  
  summarize(NumberOfCustomers = n())
```

and when predicting the testset we get the following accuracy:

```
set.seed(9)  
idx <- sample(nrow(kmeansSegments), size = nrow(test_set), replace=TRUE, prob = kmeansSegments$NumberOfCustomers)  
test_pred <- kmeansSegments[idx, ]$kmeansSegment  
accuracy <- mean(test_pred == test_set$kmeansSegment)  
accuracy_random_onKmeans = sprintf("%.2f", accuracy)  
accuracy_random_onKmeans
```

```
## [1] "0.35"
```

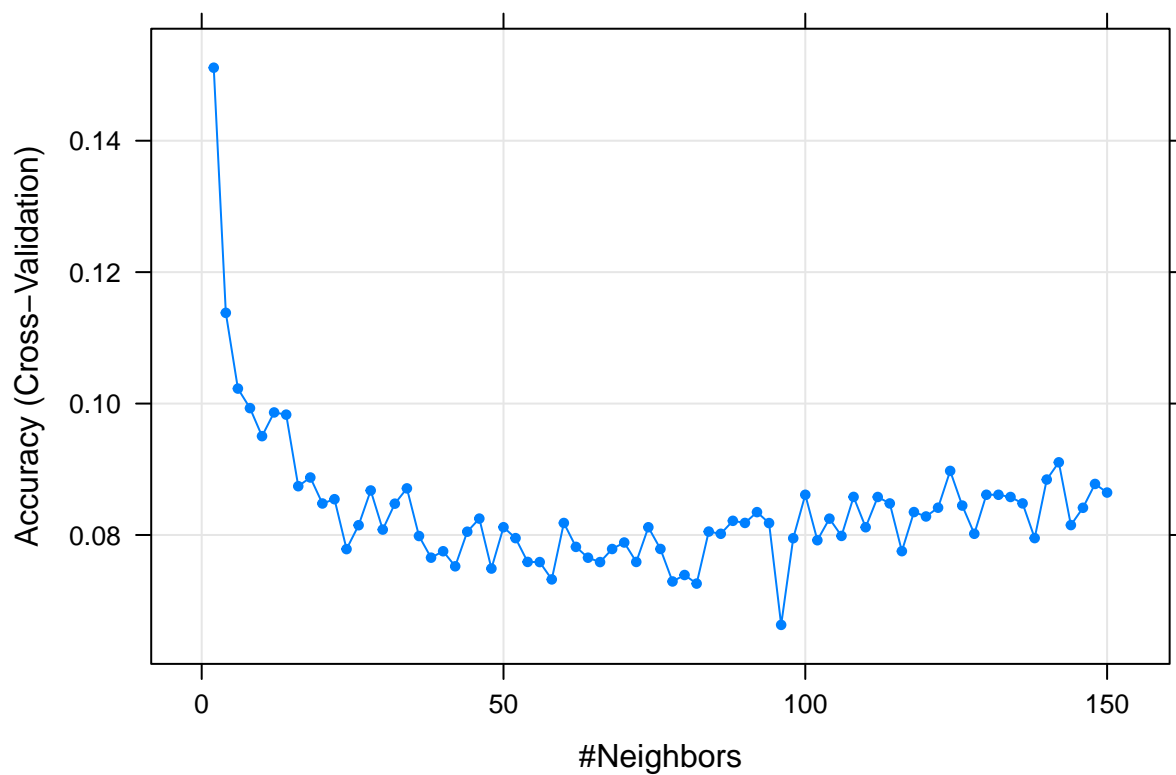
## Train and test K-nearest Neighbor Algorithm for Putler Segments

### Training

I use the caret package that takes care of cross validation and running on different values of  $k$ , the number of neighbors to look at, to find the best value. First I try to find the best  $k$  for Putler segments using 5-fold cross validation:

```
set.seed(9)
trained_knn_onPutler <- train(putlerSegment ~ .,
                              method = "knn",
                              data = select(train_set, -kmeansSegment),
                              trControl = trainControl(method = "cv", number = 5),
                              tuneGrid = data.frame(k = seq(2, 150, 2)))
```

```
plot(trained_knn_onPutler, pch = 20)
```



The algorithm finds the best  $k = 2$  which I find is surprisingly low.

### Results

Using  $k = 2$  I now find the accuracy when trying to predict the testset:

```
test_pred <- predict(trained_knn_onPutler, newdata = test_set)
accuracy <- mean(test_pred == test_set$putlerSegment)
accuracy_knn_onPutler = sprintf("%.2f", accuracy)
accuracy_knn_onPutler
```

```
## [1] "0.14"
```

Looking at the Confusion Matrix it is seen that the method predicts all segments and in different frequency as expected and there seems to be no problems with class imbalance:

```
confusionMatrix(test_pred, test_set$putlerSegment)$table
```

```
##
##           Reference
## Prediction  AboutToSleep AtRisk CantLoseThem Champions Hibernating
## AboutToSleep      18      15          0         12        14
## AtRisk            16      24          0         15        14
## CantLoseThem       0       1          0          0         0
## Champions         14      15          0         14         7
## Hibernating        8      17          0         13         4
## Lost              16       8          0          9         8
## LoyalCustomers     10      19          0         20        11
## NeedingAttention    5      10          0          9         4
## PotentialLoyalist  21      24          0         14        14
## Promising          8       2          0          2         5
## RecentCustomers    2       1          0          4         2
##
##           Reference
## Prediction  Lost LoyalCustomers NeedingAttention PotentialLoyalist
## AboutToSleep      36          33          3          26
## AtRisk            38          36          6          23
## CantLoseThem       2           0          0           0
## Champions         27          25          8          21
## Hibernating        13          28          2          14
## Lost              44          25          6          23
## LoyalCustomers     43          46         12          29
## NeedingAttention   19          17          2          10
## PotentialLoyalist  49          29          7          25
## Promising          21           9          2          11
## RecentCustomers    5           3          2           3
##
##           Reference
## Prediction  Promising RecentCustomers
## AboutToSleep      11          6
## AtRisk            4           3
## CantLoseThem       0           0
## Champions          4           1
## Hibernating        2           1
## Lost               5           5
## LoyalCustomers     6           2
## NeedingAttention    1           1
## PotentialLoyalist   7           5
## Promising           5           0
## RecentCustomers     2           0
```


## Train and Test K-nearest Neighbor Algorithm for Kmeans Segments

### Training

Again I first try to find the best k for Kmeans segments using 5-fold cross validation:

```
set.seed(9)
trained_knn_onKmeans <- train(kmeansSegment ~ .,
                              method = "knn",
                              data = select(train_set, -putlerSegment),
                              trControl = trainControl(method = "cv", number = 5),
                              tuneGrid = data.frame(k = seq(2, 150, 2)))
```

```
plot(trained_knn_onKmeans, pch = 20)
```



As for the Putler segments the best k is surprisingly low at k = 2.

### Results

Using k = 2 I now find the accuracy when trying to predict the testset:

```
test_pred <- predict(trained_knn_onKmeans, newdata = test_set)
accuracy <- mean(test_pred == test_set$kmeansSegment)
accuracy_knn_onKmeans = sprintf("%.2f", accuracy)
accuracy_knn_onKmeans
```

```
## [1] "0.31"
```

Looking at the Confusion Matrix it is seen that the method predicts all segments and in different frequency as expected and there seems to be no problems with class imbalance:

```
confusionMatrix(test_pred, test_set$kmeansSegment)$table
```

```
##           Reference
## Prediction  1   2   3   4   5   6   7
##           1  20  82   3   0  15  32   0
##           2  49 276   4   0  78  75   1
##           3   0   9   0   0   1   3   0
##           4   0   1   0   0   0   0   0
##           5  22 169   2   1  37  52   0
##           6  30 218   4   0  47  70   1
##           7   0   1   0   0   0   0   0
```

## Train and Test Random Forest Algorithm for Putler Segments

### Training


A Random Forest works by building many different Decision Trees during training. The predictions are made by majority vote by the decision trees.

The Random Forest algorithm has several parameters - the one that is said to have most influence on the result is `mtry` that is the number of variables randomly sampled as candidates at each split of decision tree.

Using 5-fold cross validation and the default of 500 decision trees the training is:

```
set.seed(9)
trained_Forest_onPutler <- train(putlerSegment ~ .,
                                method = "rf",
                                data = select(train_set, -kmeansSegment),
                                trControl = trainControl(method = "cv", number = 5),
                                tuneGrid = data.frame(mtry = seq(1,15,1)))
```

```
plot(trained_Forest_onPutler, pch = 20)
```



## Results

The algorithm found the best `mtry` = 15 but it is at the end of my search interval. From different other peoples experience I expected `mtry` to be lower than 15 but I was wrong. Because the training is very slow and because the curve seems to flatten at 15 I choose to keep `mtry` = 15. Using this value I now find the accuracy when trying to predict the testset:

```
test_pred <- predict(trained_Forest_onPutler, newdata = test_set)
accuracy <- mean(test_pred == test_set$putlerSegment)
accuracy_Forest_onPutler = sprintf("%.2f", accuracy)
accuracy_Forest_onPutler
```

```
## [1] "0.18"
```

Looking at the Confusion Matrix it is seen that the method predicts all segments and in different frequency as expected and there seems to be no problems with class imbalance:

```
confusionMatrix(test_pred, test_set$putlerSegment)$table
```

```
##
## Prediction      Reference
## AboutToSleep    AboutToSleep AtRisk CantLoseThem Champions Hibernating
## AboutToSleep          9      16          0         12          9
## AtRisk                9      17          0         13         10
## CantLoseThem         0       0          0          0          1
```



##	Champions	15	20	0	23	8
##	Hibernating	3	9	0	6	2
##	Lost	25	21	0	19	18
##	LoyalCustomers	13	25	0	22	15
##	NeedingAttention	10	8	0	7	3
##	PotentialLoyalist	12	11	0	7	11
##	Promising	15	6	0	3	3
##	RecentCustomers	7	3	0	0	3
##		Reference				
##	Prediction	Lost	LoyalCustomers	NeedingAttention	PotentialLoyalist	
##	AboutToSleep	24	27	6	28	
##	AtRisk	17	31	4	12	
##	CantLoseThem	1	1	0	0	
##	Champions	34	25	7	19	
##	Hibernating	7	19	0	8	
##	Lost	91	38	11	35	
##	LoyalCustomers	31	58	9	31	
##	NeedingAttention	21	10	2	12	
##	PotentialLoyalist	35	20	7	17	
##	Promising	25	15	2	14	
##	RecentCustomers	11	7	2	9	
##		Reference				
##	Prediction	Promising	RecentCustomers			
##	AboutToSleep	10	4			
##	AtRisk	2	1			
##	CantLoseThem	0	0			
##	Champions	1	0			
##	Hibernating	0	1			
##	Lost	7	6			
##	LoyalCustomers	4	3			
##	NeedingAttention	5	0			
##	PotentialLoyalist	5	4			
##	Promising	9	2			
##	RecentCustomers	4	3			


## Train and test Random Forest Algorithm for Kmeans Segments

### Training

For the Kmeans segments I train:

```
set.seed(9)
trained_Forest_onKmeans <- train(kmeansSegment ~ .,
                                method = "rf",
                                data = select(train_set, -putlerSegment),
                                trControl = trainControl(method = "cv", number = 5),
                                tuneGrid = data.frame(mtry = seq(1,15,1)))

plot(trained_Forest_onKmeans, pch = 20)
```



## Results

Again the algorithm found the best `mtry` = 15 and as before I choose to keep `mtry` = 15. Using this value I now find the accuracy when trying to predict the testset:

```
test_pred <- predict(trained_Forest_onKmeans, newdata = test_set)
accuracy <- mean(test_pred == test_set$kmeansSegment)
accuracy_Forest_onKmeans = sprintf("%.2f", accuracy)
accuracy_Forest_onKmeans
```

```
## [1] "0.39"
```

Looking at the Confusion Matrix it is seen that the method predicts all segments and in different frequency as expected and there seems to be no problems with class imbalance:

```
confusionMatrix(test_pred, test_set$kmeansSegment)$table
```

```
##           Reference
## Prediction  1  2  3  4  5  6  7
##           1 21 83  2  0 15 19  0
##           2 66 417  5  0 89 124  1
##           3  0 12  0  0  2  4  0
##           4  0  1  0  0  0  0  0
##           5 20 136  4  1 40 52  0
```

##	6	14	106	2	0	32	33	1
##	7	0	1	0	0	0	0	0

## Collected Results of Predicting

Method	Putler Segments Accuracy	Kmeans Segments Accuracy
Random	0.13	0.35
K-nearest neighbor	0.14	0.31
Random Forest	0.18	0.39

## Conclusion

### Customer Segments

The traditional, fixed RFM analysis that I used defines eleven different customer segments. The Machine Learning generated customer segments turn out to have only seven customer segments. The analysis shows that the two sets of customer segments splits the customers in very different ways.

The machine learning method created a reasonable sized set of segments that when interpreted seems to make good business sense. It is, however, impossible to say which set of the segment sets that are best because it of course depends on the intended use. If we want a set of segments that is good at describing the states a customer goes through during time, it might be possible to make an analysis of the development over time of customers. In this report I assumed that the customers were close to static over the year of the dataset. I believe that finding dynamic trends or customer lifecycles requires a dataset that covers a longer period of at least several years for a business like the one, I analysed in this report.

### Predicting Customer Segment

Using the K-nearest neighbor algorithm to predict the future customer segment from customers first order and their country did worse than random for the Kmeans segment! However, the Random Forest algorithm predicted the future customer segment from customers first order and their country a little better than random for both segment sets. Even though better than random, the accuracy is still low for both segment sets. Predicting using Random Forest might, though, be a beginning of helping the business to a finer segmentation of new customers.

It is clearly easier to predict the Kmeans segments - this is quite naturally because of the fewer segments and the very high number of customers in a single segment.

In this report I did not use the information of the content of the first order. It would be interesting to investigate if using this data, or perhaps even making a Market Basket Analysis, could help increase the accuracy of the predicting. Also I should refine the Random Forest analyses to be sure to find the really best value of `mtry`.